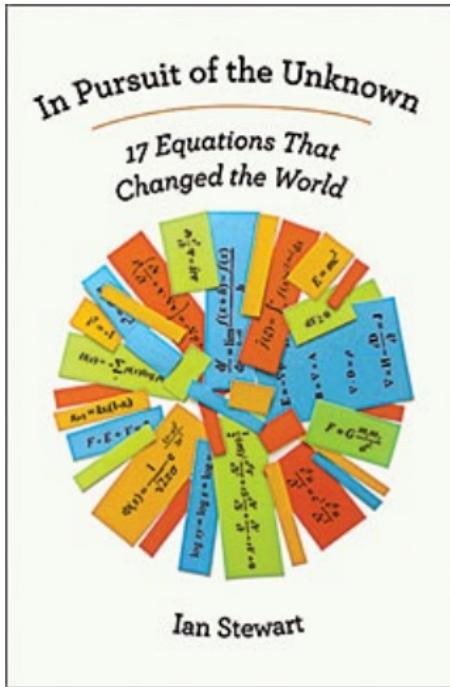


Image Features

So far

- In images (practically) the lowest level of organization (Pixel)
- Filtering (Extract structure from a collection of pixels)
- Convolution (a key operations) (Implemented in working systems through Fourier transform)

Seventeen Equations.....



Seventeen Equations.....

17 Equations That Changed the World by Ian Stewart

- | | | | |
|-----|-------------------------------|---|----------------------------|
| 1. | Pythagoras's Theorem | $a^2 + b^2 = c^2$ | Pythagoras, 530 BC |
| 2. | Logarithms | $\log xy = \log x + \log y$ | John Napier, 1610 |
| 3. | Calculus | $\frac{df}{dt} = \lim_{h \rightarrow 0} \frac{f(t+h) - f(t)}{h}$ | Newton, 1668 |
| 4. | Law of Gravity | $F = G \frac{m_1 m_2}{r^2}$ | Newton, 1687 |
| 5. | The Square Root of Minus One | $i^2 = -1$ | Euler, 1750 |
| 6. | Euler's Formula for Polyhedra | $V - E + F = 2$ | Euler, 1751 |
| 7. | Normal Distribution | $\Phi(x) = \frac{1}{\sqrt{2\pi\rho}} e^{-\frac{(x-\mu)^2}{2\rho^2}}$ | C.F. Gauss, 1810 |
| 8. | Wave Equation | $\frac{\partial^2 u}{\partial t^2} = c^2 \frac{\partial^2 u}{\partial x^2}$ | J. d'Almbert, 1746 |
| 9. | Fourier Transform | $f(\omega) = \int_{-\infty}^{\infty} f(x) e^{-2\pi i \omega x} dx$ | J. Fourier, 1822 |
| 10. | Navier-Stokes Equation | $\rho \left(\frac{\partial \mathbf{v}}{\partial t} + \mathbf{v} \cdot \nabla \mathbf{v} \right) = -\nabla p + \nabla \cdot \mathbf{T} + \mathbf{f}$ | C. Navier, G. Stokes, 1845 |
| 11. | Maxwell's Equations | $\begin{aligned} \nabla \cdot \mathbf{E} &= \frac{\rho}{\epsilon_0} & \nabla \cdot \mathbf{H} &= 0 \\ \nabla \times \mathbf{E} &= -\frac{1}{c} \frac{\partial \mathbf{H}}{\partial t} & \nabla \times \mathbf{H} &= \frac{1}{c} \frac{\partial \mathbf{E}}{\partial t} \end{aligned}$ | J.C. Maxwell, 1865 |
| 12. | Second Law of Thermodynamics | $dS \geq 0$ | L. Boltzmann, 1874 |
| 13. | Relativity | $E = mc^2$ | Einstein, 1905 |
| 14. | Schrodinger's Equation | $i\hbar \frac{\partial}{\partial t} \Psi = H\Psi$ | E. Schrodinger, 1927 |
| 15. | Information Theory | $H = - \sum p(x) \log p(x)$ | C. Shannon, 1949 |
| 16. | Chaos Theory | $x_{t+1} = kx_t(1-x_t)$ | Robert May, 1975 |
| 17. | Black-Scholes Equation | $\frac{1}{2} \sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} + \frac{\partial V}{\partial t} - rV = 0$ | F. Black, M. Scholes, 1990 |

So far

- In images (practically) the lowest level of organization (Pixel)
- Filtering (Extract structure from a collection of pixels)
- Convolution (a key operations) (Implemented in working systems through Fourier transform)
- Smoothening (Gaussian)
- Edges (Good Abstraction)
- Role of Derivatives
- Scale of the images (Downsample Repeatedly Pyramid)
- Subsampling (Nyquist Rate)
- Upsampling (Interpolation)
- Interpolation also implemented through convolutions

Image Features

- What skyline is this?



Image Features

- What skyline is this?



Image Features

- What skyline is this?

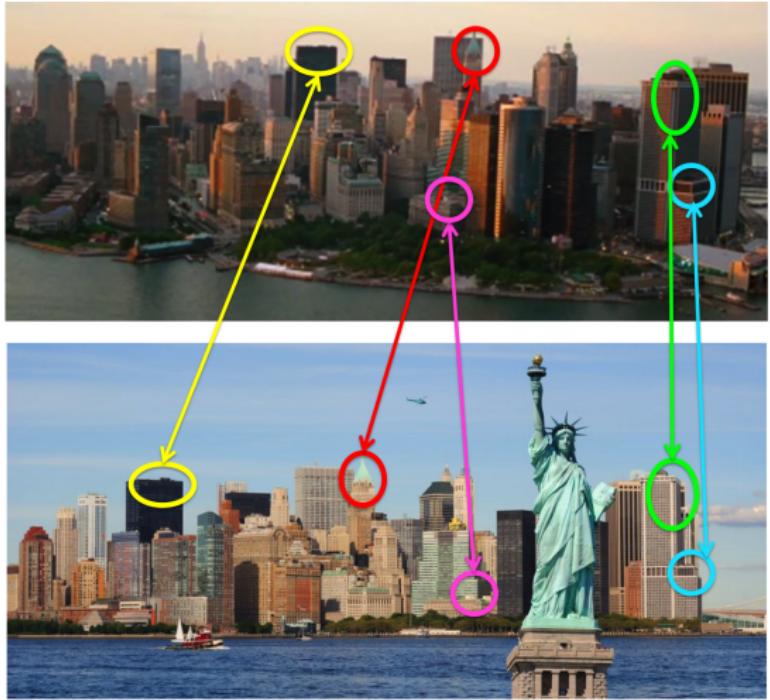


Image Features

- What skyline is this?

We matched in:

- Distinctive locations:
keypoints
- Distinctive features:
descriptors

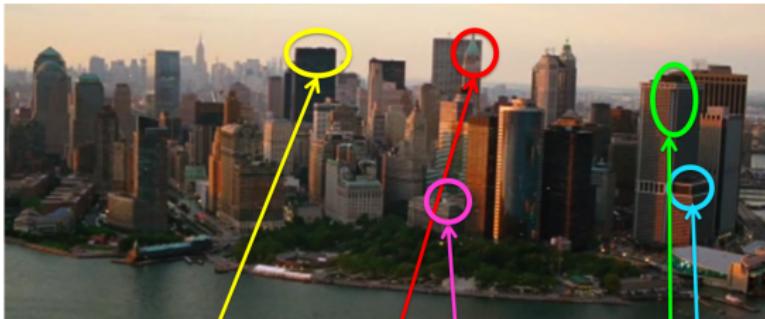


Image Features

- How could we tell which type of scene this is from a movie?



What kind of scene is behind the actors?
Kitchen? Bedroom? Street? Bar?

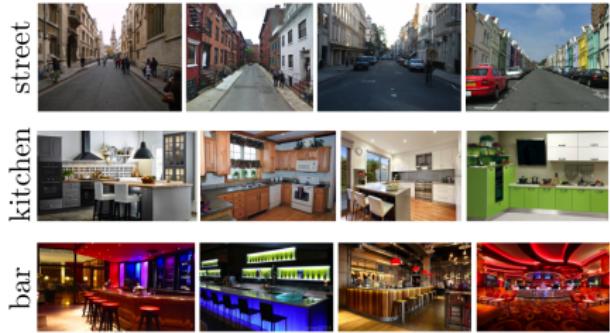


Image Features

- How could we tell which type of scene this is from a movie?

We matched:

- **Globally** – one descriptor for full image (?)
- More complex descriptor: color, gradients, “deep” features (learned), etc



What kind of scene is behind the actors?

Kitchen? Bedroom? Street? Bar?

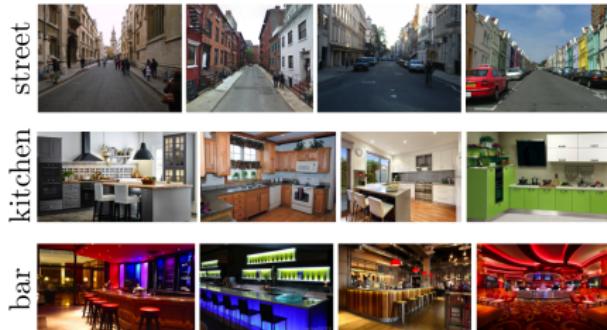


Image Features

- How would we solve this?



Are these two cups of the same type?

Image Features

- How would we solve this?

We matched:

- One descriptor for full **patch**
- Descriptor can be simple, e.g. **color**



Are these two cups of the same type?

Image Features

- How would we solve this?



Where can I find this pattern? ➔ **LAKE BELL**

Image Features

- How would we solve this?

We matched:

- At each location
- Compared pixel values



Where can I find this pattern? ➔ **LAKE BELL**

Image Features

- How would we solve this?



Where can I find this pattern?>



Image Features

- How would we solve this?

We matched:

- Distinctive locations
- Distinctive features
- Affine invariant



Where can I find this pattern? ➔

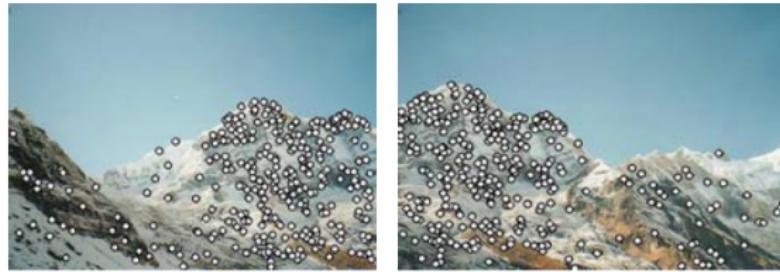
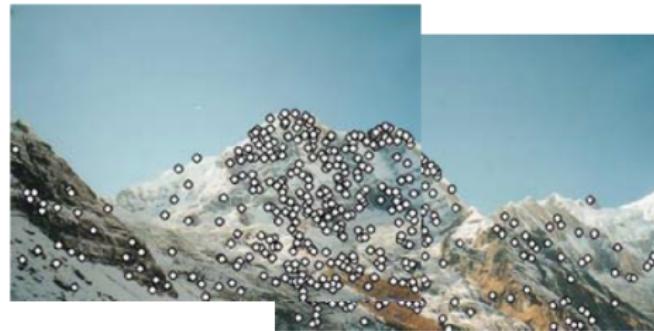


Image Features

- How would we solve this?

Image Features:
Interest Point (Keypoint) Detection

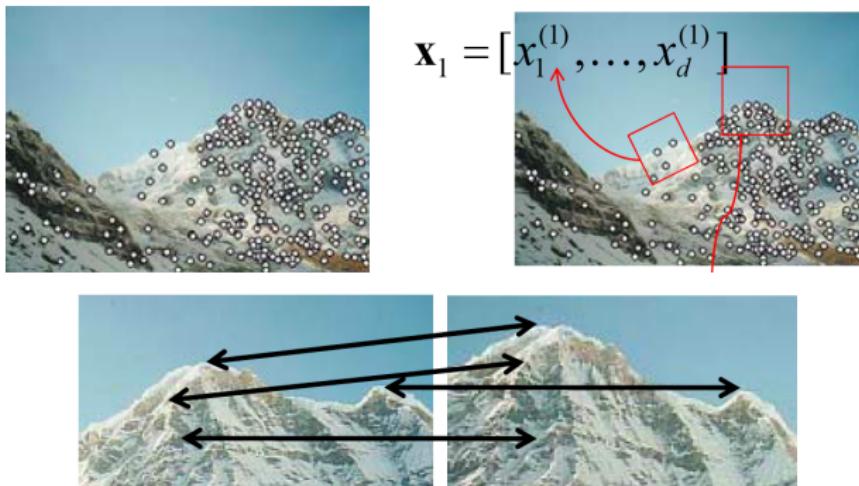
Application Example: Image Stitching



[Source: K. Grauman]

Local Features

- **Detection:** Identify the interest points.
- **Description:** Extract **feature vector** descriptor around each interest point.
- **Matching:** Determine correspondence between descriptors in two views.



[Source: K. Grauman]

Goal: Repeatability of the Interest Point Operator

- Our goal is to detect (at least some of) the same points in both images
- We have to be able to run the detection procedure independently per image
- We need to generate enough points to increase our chances of detecting matching points
- We shouldn't generate too many or our matching algorithm will be too slow



Figure: Too few keypoints → little chance to find the true matches

[Source: K. Grauman, slide credit: R. Urtasun]

What Points to Choose?

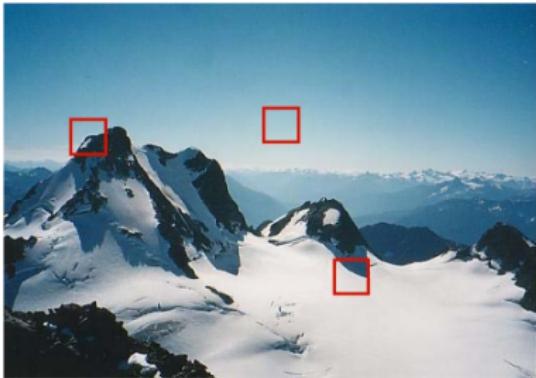


[Source: K. Grauman]

What Points to Choose for matching?



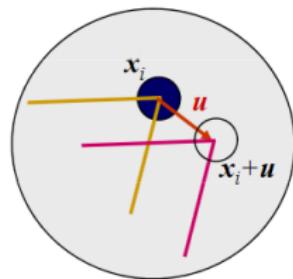
What Points to Choose for matching?



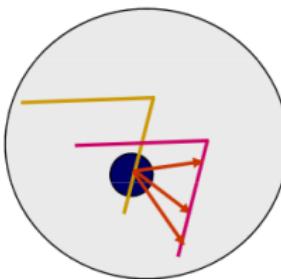
- Textureless patches are nearly impossible to localize.
- Patches with large contrast changes (gradients) are easier to localize.
- But straight line segments cannot be localized on lines segments with the same orientation (aperture problem)
- Gradients in at least two different orientations are easiest, e.g., **corners!**

[Adopted from: Szelski (Book)]

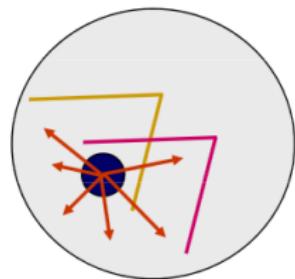
Aperture Problem



(a)



(b)



(c)

- “Corner-like” patch can be reliably matched
- A straight line patch can have multiple matches (Aperture Problem)
- Zero texture, useless, can have infinite matches

[Adopted from: Szelski (Book)]

Interest Points: Corners

- How can we find corners in an image?



Interest Points: Corners

- We should easily recognize the point by looking through a small window.
- Shifting a window in any direction should give a large change in intensity.

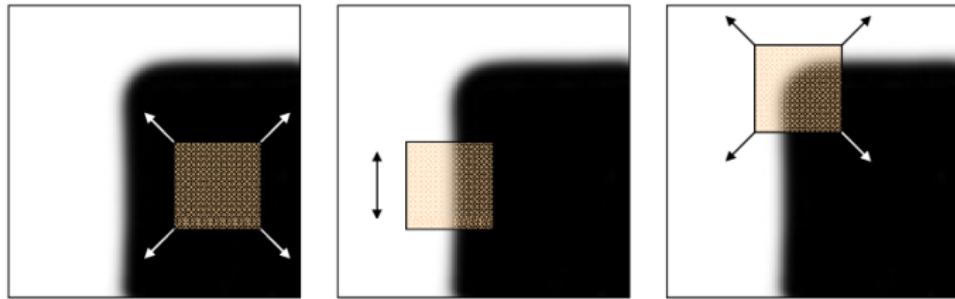


Figure: (left) flat region: no change in all directions, (center) edge: no change along the edge direction, (right) corner: significant change in all directions

[Source: Alyosha Efros, Darya Frolova, Denis Simakov]

Interest Points: Corners

- Harris Corner Detector: Idea



$\sum I_x^2$ is large

$\sum I_y^2$ is large

⇒ Corner!

Interest Points: Corners

- Harris Corner Detector: Idea



$$\begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix}$$

\implies eigenvalues!

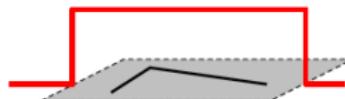
Interest Points: Corners

- Compare two image patches using (weighted) summed square difference
- Measures change in appearance of window $w(x, y)$ for the shift

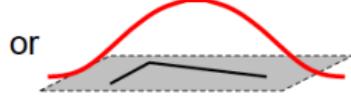
$$E_{\text{WSSD}}(u, v) = \sum_x \sum_y w(x, y)[I(x + u, y + v) - I(x, y)]^2$$

↑
window function ↑
shifted intensity ↑
intensity

Window function $w(x, y) =$



1 in window, 0 outside



Gaussian

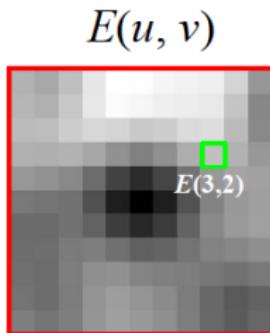
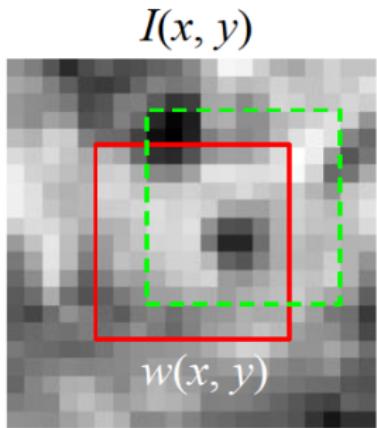
[Source: J. Hays]

Interest Points: Corners

- Compare two image patches using (weighted) summed square difference
- Measures change in appearance of window $w(x, y)$ for the shift

$$E_{\text{WSSD}}(u, v) = \sum_x \sum_y w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

↑
window function ↑
shifted intensity ↑
intensity



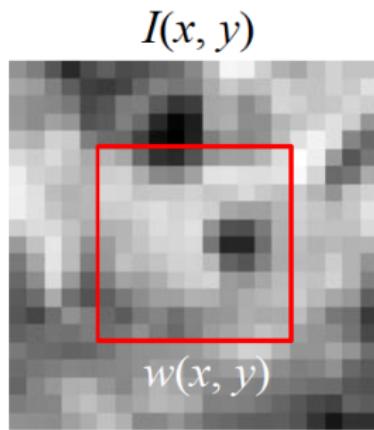
[Source: J. Hays]

Interest Points: Corners

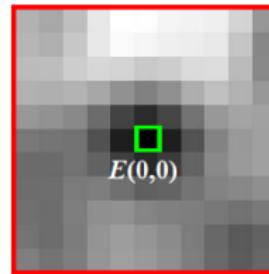
- Compare two image patches using (weighted) summed square difference
- Measures change in appearance of window $w(x, y)$ for the shift

$$E_{\text{WSSD}}(u, v) = \sum_x \sum_y w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

↑
window function ↑
shifted intensity ↑
intensity



$E(u, v)$

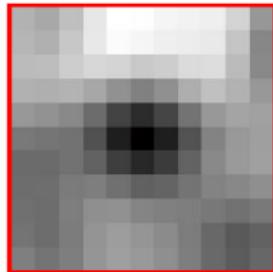


[Source: J. Hays]

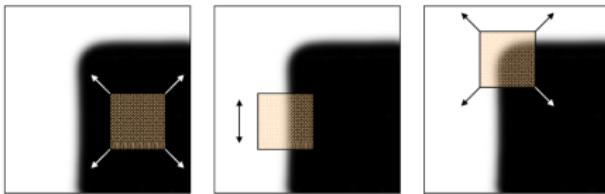
Interest Points: Corners

- Let's look at E_{WSSD}
- We want to find out how this function behaves for small shifts

$$E(u, v)$$



- Remember our goal to detect corners:



Interest Points: Corners

- Using a simple first-order Taylor Series expansion about x, y :

$$I(x + u, y + v) \approx I(x, y) + u \cdot \frac{\partial I}{\partial x}(x, y) + v \cdot \frac{\partial I}{\partial y}(x, y)$$

- Using a series of polynomials to approximate I , more info on Taylor Series [here](#)
- And plugging it in our expression for E_{WSSD} :

$$\begin{aligned} E_{WSSD}(u, v) &= \sum_x \sum_y w(x, y) \left(I(x + u, y + v) - I(x, y) \right)^2 \\ &\approx \sum_x \sum_y w(x, y) \left(I(x, y) + u \cdot I_x + v \cdot I_y - I(x, y) \right)^2 \\ &= \sum_x \sum_y w(x, y) \left(u^2 I_x^2 + 2u \cdot v \cdot I_x \cdot I_y + v^2 I_y^2 \right) \\ &= \sum_x \sum_y w(x, y) \cdot \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

Interest Points: Corners

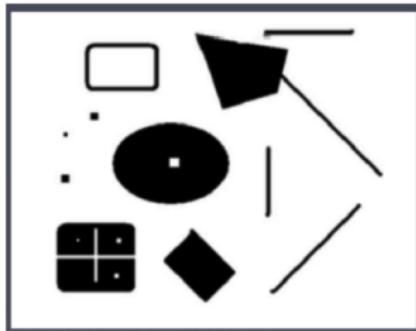
- Since (u, v) doesn't depend on (x, y) we can rewrite it slightly:

$$\begin{aligned} E_{\text{WSSD}}(u, v) &= \sum_x \sum_y w(x, y) \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \\ &= \begin{bmatrix} u & v \end{bmatrix} \underbrace{\left(\sum_x \sum_y w(x, y) \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix} \right)}_{\text{Let's denote this with } M} \begin{bmatrix} u \\ v \end{bmatrix} \\ &= \begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

- M is a 2×2 *second moment matrix* computed from image gradients:

$$M = \sum_x \sum_y w(x, y) \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix}$$

How Do I Compute M ?

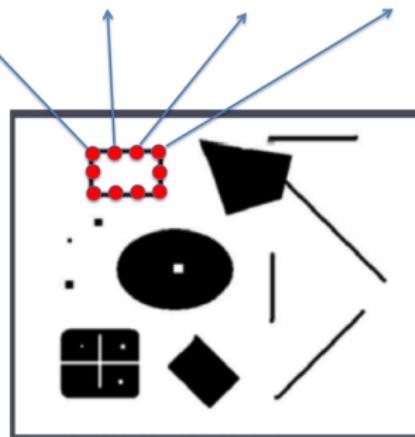


image

- Let's say I have this image

How Do I Compute M ?

$$M = ? \quad M = ? \quad M = ? \quad M = ?$$

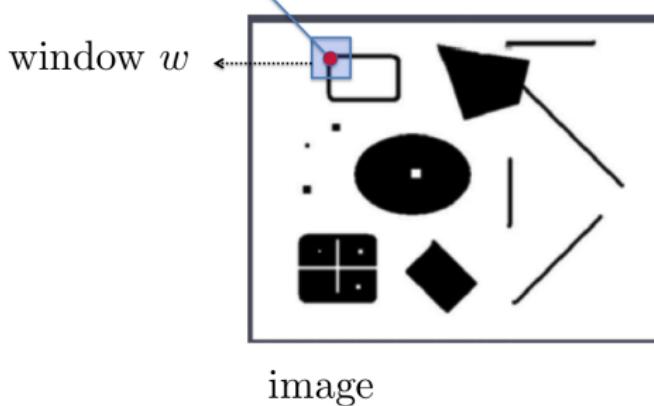


image

- Let's say I have this image
- I need to compute a 2×2 second moment matrix in each image location

How Do I Compute M ?

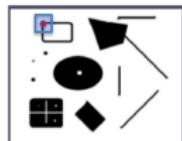
$$M = \sum_x \sum_y w(x, y) \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix}$$



- Let's say I have this image
- I need to compute a 2×2 second moment matrix in each image location
- In a particular location I need to compute M as a weighted average of gradients in a window

How Do I Compute M ?

$$M = \sum_x \sum_y w(x, y) \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix}$$



image



$$I_x = \frac{\partial I}{\partial x}$$



$$I_y = \frac{\partial I}{\partial y}$$



$$I_x \cdot I_y$$

- Let's say I have this image
- I need to compute a 2×2 second moment matrix in each image location
- In a particular location I need to compute M as a weighted average of gradients in a window

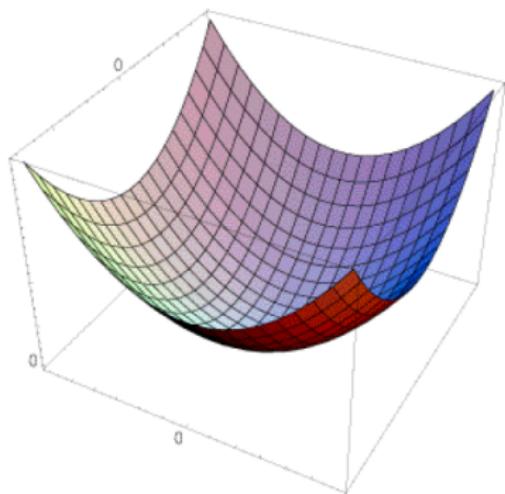
I can do this efficiently by computing three matrices, I_x^2 , I_y^2 and $I_x \cdot I_y$, and convolving each one with a filter, e.g. a box or Gaussian filter

Interest Points: Corners

- We now have M computed in each image location
- Our E_{WSSD} is a **quadratic function** where M implies its shape

$$E_{\text{WSSD}}(u, v) = [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_x \sum_y w(x, y) \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix}$$



[Source: J. Hays]

Interest Points: Corners

- Let's take a horizontal "slice" of $E_{WSSD}(u, v)$:

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

- This is the equation of an ellipse

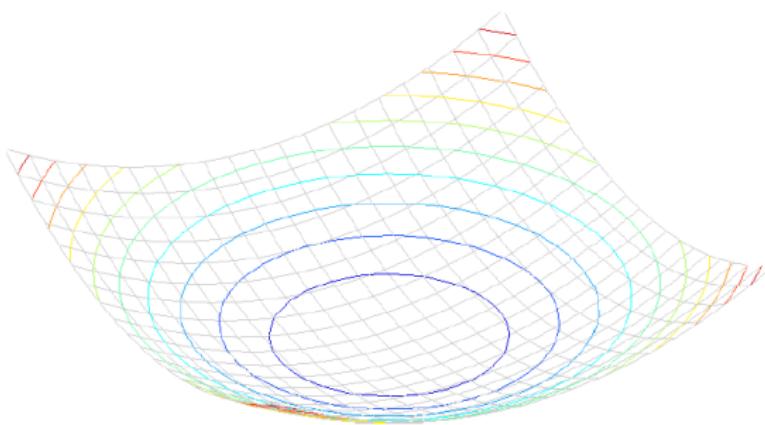


Figure: Different ellipses obtain by different horizontal "slices"

Interest Points: Corners

- Let's take a horizontal "slice" of $E_{WSSD}(u, v)$:

$$[u \quad v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

- This is the equation of an ellipse

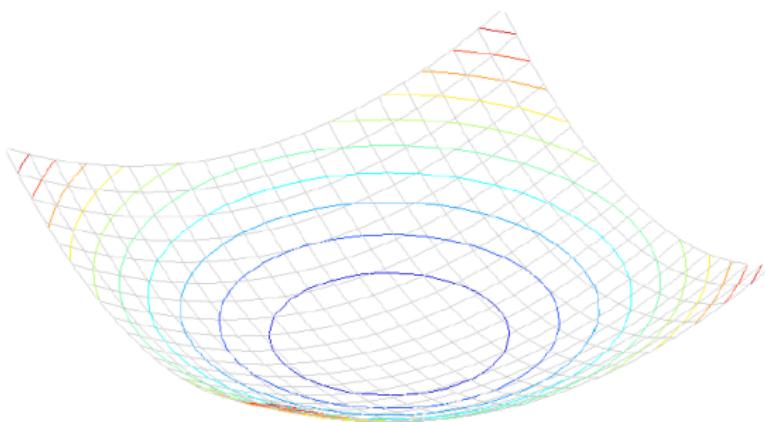


Figure: Different ellipses obtain by different horizontal "slices"

Interest Points: Corners

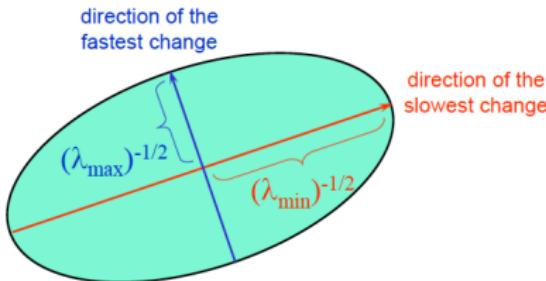
- Our matrix M is symmetric:

$$M = \sum_x \sum_y w(x, y) \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix}$$

- And thus we can diagonalize it (in Matlab: $[V, D] = \text{EIG}(M)$):

$$M = V \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} V^{-1}$$

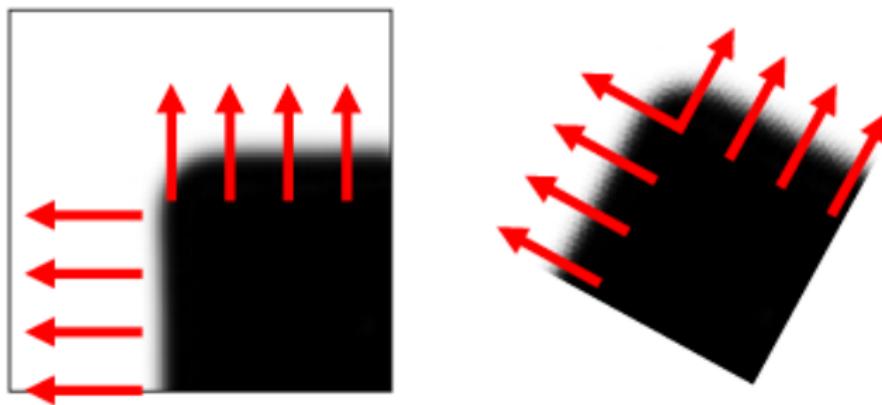
- Columns of V are major and minor axes of ellipse, the lengths of the radii proportional to $\lambda^{-1/2}$



[Source: J. Hays]

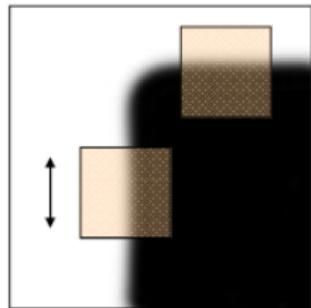
Interest Points: Corners

- The eigenvalues of M (λ_1, λ_2) reveal the amount of intensity change in the two principal orthogonal gradient directions in the window



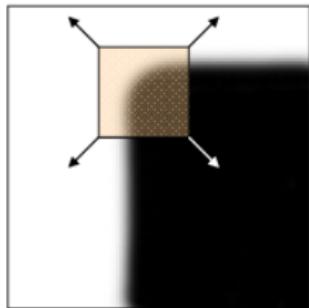
[Source: R. Szeliski, slide credit: R. Urtasun]

Interest Points: Corners



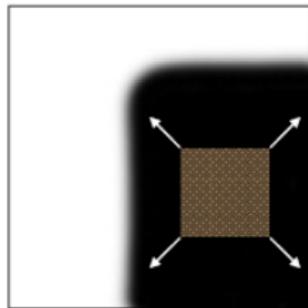
“edge”:

$$\begin{aligned}\lambda_1 &>> \lambda_2 \\ \lambda_2 &>> \lambda_1\end{aligned}$$



“corner”:

$$\begin{aligned}\lambda_1 \text{ and } \lambda_2 \text{ are large,} \\ \lambda_1 \sim \lambda_2;\end{aligned}$$



“flat” region

$$\begin{aligned}\lambda_1 \text{ and } \lambda_2 \text{ are} \\ \text{small;}\end{aligned}$$

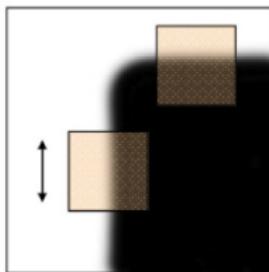
[Source: K. Grauman, slide credit: R. Urtasun]

Interest Points: Criteria to Find Corners

- Harris and Stephens, '88, is rotationally invariant and downweights edge-like features where $\lambda_1 \gg \lambda_0$

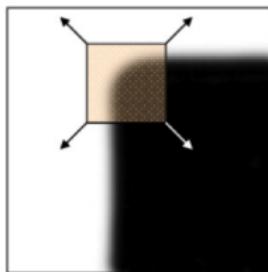
$$R = \lambda_0\lambda_1 - \alpha(\lambda_0 + \lambda_1)^2 = \det(M) - \alpha \cdot \text{trace}(M)^2$$

- Why go via det and trace and not use a criteria with λ ?
- α a constant (0.04 to 0.06)



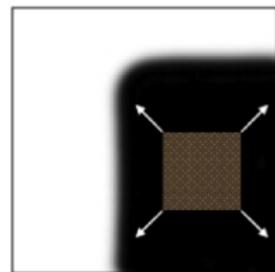
“edge”:

$$R < 0$$



“corner”:

$$R > 0$$



“flat” region

$$|R| \text{ small}$$

- The corresponding detector is called **Harris corner detector**

Interest Points: Criteria to Find Corners

- Harris and Stephens, 88 is rotationally invariant and downweights edge-like features where $\lambda_1 \gg \lambda_0$

$$R = \lambda_0\lambda_1 - \alpha(\lambda_0 + \lambda_1)^2 = \det(M) - \alpha\text{trace}(M)^2$$

- Shi and Tomasi, 94 proposed the smallest eigenvalue of \mathbf{A} , i.e., $\lambda_0^{-1/2}$.
- Triggs, 04 suggested

$$\lambda_0 - \alpha\lambda_1$$

also reduces the response at 1D edges, where aliasing errors sometimes inflate the smaller eigenvalue

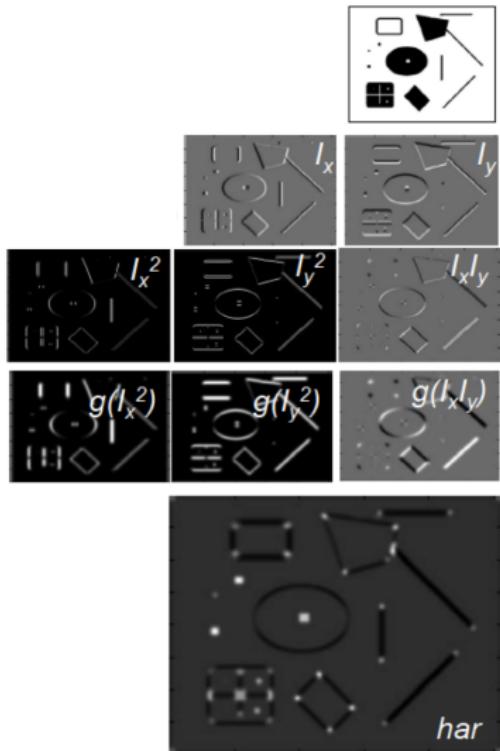
- Brown et al, 05 use the harmonic mean

$$\frac{\det(\mathbf{A})}{\text{trace}(\mathbf{A})} = \frac{\lambda_0\lambda_1}{\lambda_0 + \lambda_1}$$

[Source Mubarak Shah, Szelski]

Harris Corner detector

- ① Compute gradients I_x and I_y
- ② Compute I_x^2 , I_y^2 , $I_x \cdot I_y$
- ③ Average (Gaussian) → gives M per voxel
- ④ Compute
 $R = \det(M) - \alpha \text{trace}(M)^2$ for each image window (*cornerness score*)
- ⑤ Find points with large R ($R >$ threshold).
- ⑥ Take only points of local maxima,
i.e., perform non-maximum suppression

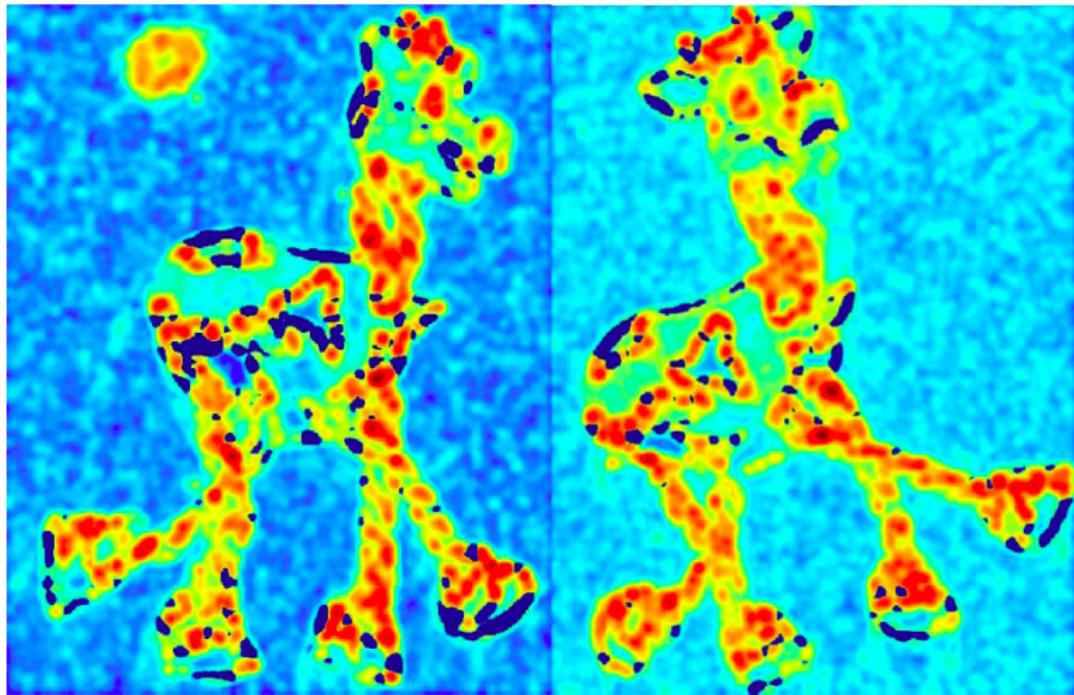


Example



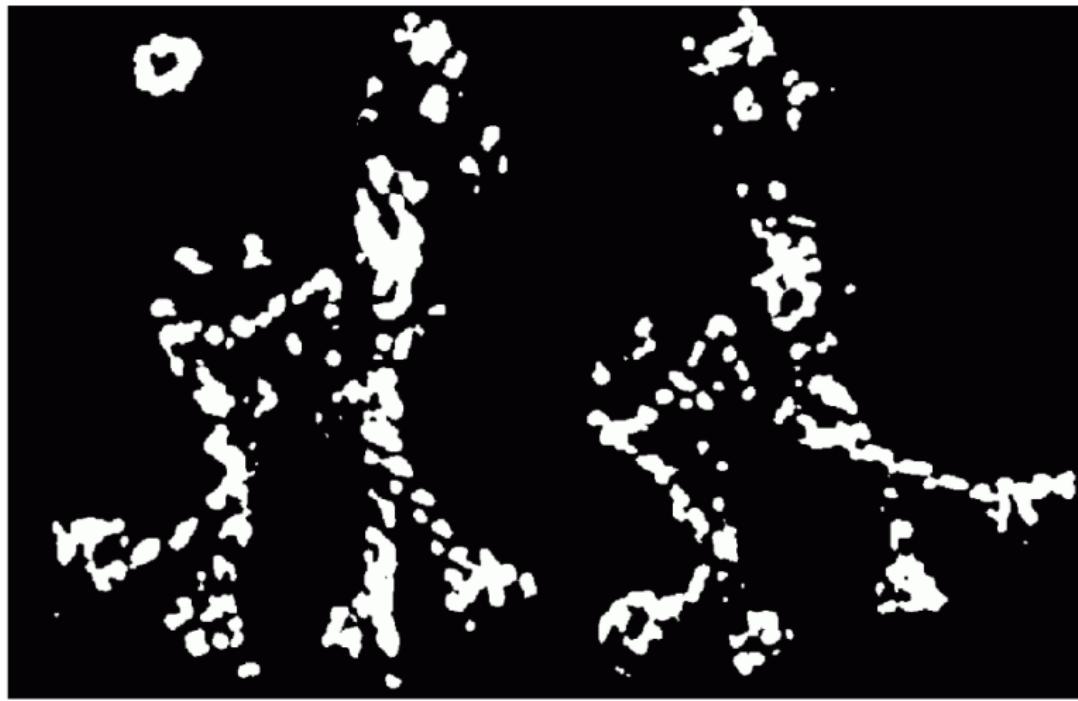
[Source: K. Grauman]

1) Compute Cornerness



[Source: K. Grauman]

2) Find High Response



[Source: K. Grauman]

3) Non-maxima Suppresion



[Source: K. Grauman]

Results



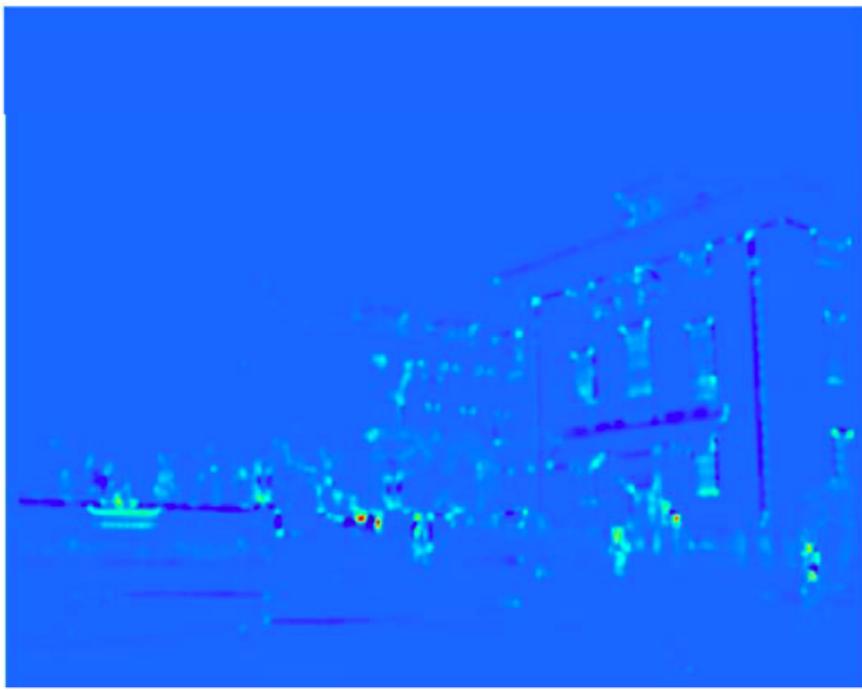
[Source: K. Grauman]

Another Example



[Source: K. Grauman]

Cornerness



[Source: K. Grauman]

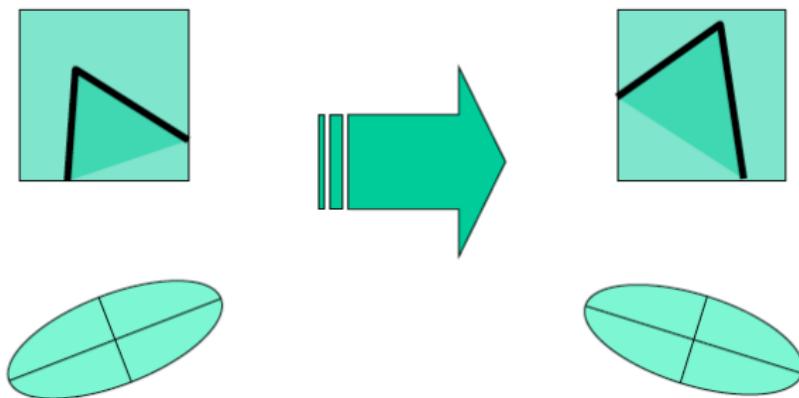
Interest Points



[Source: K. Grauman]

Properties of Harris Corner Detector

- Rotation and Shift Invariance of Corners

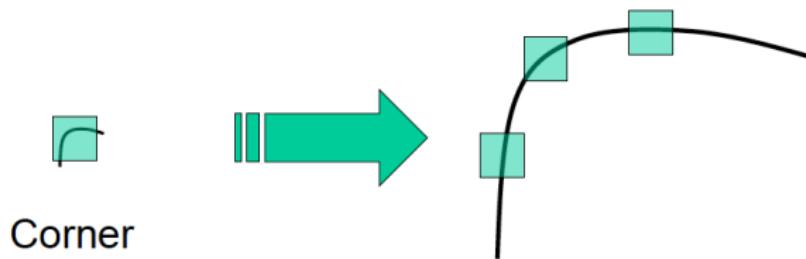


- Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same
- Harris corner detector is rotation-covariant

[Source: J. Hays]

Properties of Harris Corner Detector

- Scale?



All points will
be classified
as edges

- Corner location is **not scale invariant/covariant!**

[Source: J. Hays]

Next Time

- Can we also define keypoints that are shift, rotation and scale invariant/covariant?
- What should be our **description** around keypoint?